

Advanced Level MCQs — All 7 Modules

Panaversity Agent Factory AI Prompting 2026

Prepared by: Imran Hassan

ADVANCED LEVEL · SCENARIO-BASED · CROSS-MODULE

0

SCORE

0

ANSWERED

35

TOTAL QS

140

TOTAL MARKS

Module 1 — AI Agent Factory Foundations (Q1–Q5)

SCENARIO

A solo developer builds an AI agent for legal contract review. She uses the latest Claude model, writes detailed prompts, and encodes lawyer-reviewed rules into the system. However, there is no mechanism for a qualified human to review unusual contract clauses before the agent sends its analysis directly to clients. Three weeks later, the agent misclassifies a non-compete clause, causing client harm.

1 Which of the Four Critical Elements of a Digital FTE was absent, and why did its absence cause the failure?

A Domain Expertise — no lawyer-crafted rules were encoded into the agent's behavior

B Structured Specifications — the developer's prompts lacked markdown formatting

C Human Oversight — no qualified professional checkpoint existed before client-facing output was delivered, allowing an AI misclassification to reach the client unreviewed

D Engineering Architecture — the agent should have used MCP servers instead of direct API calls

SCENARIO

A content team maintains an AI education business built on a System of Record. They update the core module on "AI Prompting" with three new concepts discovered after the original publication. Their business delivers content through four channels: The Book, TutorClaw, Agent Factory Skillpack, and Derivative Books.

2 Which outcome correctly describes what happens after the System of Record is updated?

- A** Only TutorClaw updates automatically; the book and skill pack require manual edits
- B** All four delivery channels propagate the update simultaneously, maintaining a single verified source of truth across the entire business
- C** The System of Record must be re-validated by domain experts before any channel can propagate the update
- D** TutorClaw and Derivative Books update automatically, but the Book PDF requires a separate manual regeneration step

SCENARIO

A company builds a specialized AI agent that manages hospital bed allocation. It integrates with the hospital's Electronic Health Record system, generates real-time alerts for nursing staff, and learns optimal allocation patterns from patient flow history. The company's team consists of three domain experts and two AI engineers.

3 According to Jensen Huang's Five-Layer AI Model, at which layer does this company primarily operate?

- A** Layer 1 — Physical AI, because the agent controls real-world hospital operations involving patients
- B** Layer 3 — AI Agents, because the hospital bed agent operates autonomously within a defined domain

- C** Layer 5 — Applications Layer, because it is a domain-specific AI-powered application built on top of the underlying model and agent infrastructure
- D** Layer 2 — Foundation Models, because it relies on a pretrained model for all decision-making

SCENARIO

A traditional logistics company employs 500 customer service representatives handling shipment inquiries at a cost of \$8M annually. A competitor launches an AI-native logistics company with 12 employees and an AI agent that handles 95% of the same inquiry volume. Both companies serve an identical number of daily shipments.

4 What is the PRIMARY strategic implication of this comparison according to the AI Agent Factory framework?

- A** The AI-native company has a cost advantage only in the short term until customer service regulations catch up
- B** The AI-native company demonstrates the leverage principle — AI agents deliver high output with minimal headcount, compressing the traditional cost-to-output ratio dramatically
- C** The traditional company retains a long-term quality advantage because human agents handle nuanced cases better at every edge case
- D** Both models are economically equivalent because the 12-person AI company still needs 500 AI agent "seats" which cost the same

SCENARIO

A department head at a bank wants her 20-person team to use AI to draft regulatory reports faster, summarize meeting notes, and prepare client briefings. She is NOT interested in building autonomous agents or deploying new software systems — she just wants practical AI habits her team can adopt in their existing daily workflow starting this week.

5 Which learning mode should the AI Agent Factory curriculum recommend, and why?

- A** Mode 2 (Building Autonomous Agents) — regulatory environments require the reliability that only engineered systems provide
- B** Mode 1 (Using AI to Accelerate Personal Work) — this matches her stated goal of improving existing workflow without building or deploying new systems
- C** Both modes simultaneously — the hybrid approach ensures she doesn't miss the advanced techniques available in Mode 2
- D** Neither mode — a department head should delegate AI adoption entirely to IT rather than self-educating

Module 2 — AI Prompting in 2026 (Q6–Q10)**SCENARIO**

A strategy consultant needs to write a competitive analysis of a software company whose latest AI product launched two weeks ago. The product received significant press coverage but no academic papers have been published about it. She has a one-hour deadline to brief her client's board.

6 Which retrieval mode should she select and why?

- A** Pretrained mode — the foundation model has broad business and technology knowledge sufficient for competitive analysis of any software company
- B** Deep research mode — the product launched after the model's training cutoff; only live web access can retrieve current coverage, and deep research scans comprehensively across multiple sources
- C** Reasoning/thinking mode — extended reasoning compensates for any lack of recent training data
- D** Web search mode — a single-pass web search provides recent information and deep research would waste her limited one-hour deadline

SCENARIO

A product manager pitches a new feature idea to Claude: "We're thinking of adding a social sharing button to the checkout page. Is this a good idea?" Claude responds enthusiastically: "This is a great idea! Social proof at the checkout page can significantly reduce cart abandonment and boost conversions." The PM screenshots this response and shares it with executives as AI validation before the sprint planning meeting.

7 What has fundamentally gone wrong with how this validation was obtained?

- A** The PM should have used Web Search mode to independently validate market potential before asking Claude
- B** Claude's response reflects sycophancy — AI is trained to agree with user framing. Without a balanced rubric requesting pros, cons, risks, and numerical scoring, the response is approval bias, not genuine analysis
- C** The PM's question was too vague; Claude should have asked clarifying questions about the target market before responding
- D** The PM should have used Deep Research mode to generate an objective market report comparing competitors

SCENARIO

A full-stack developer spends three hours in a single Claude session: first debugging a React authentication component, then designing a PostgreSQL database schema, and finally writing OpenAPI documentation for three endpoints. By the third task, Claude starts calling database tables by the React component names and adds authentication logic into the database schema discussion — mixing context from all three tasks.

8 What is the BEST preventive habit to avoid this outcome in future sessions?

- A** Start each task by pasting the most relevant code file at the top of the existing conversation to reset Claude's attention

- B** Use Reasoning/Thinking mode for each task switch — extended reasoning resets Claude's analytical frame between topics
- C** Start a fresh conversation when switching to a substantially different task — context contamination from earlier turns progressively degrades output quality
- D** Summarize each completed task as a bulleted list before moving to the next, so Claude can distinguish each context

SCENARIO

A UX designer at a mid-sized startup wants to create 50 variations of a landing page hero section to A/B test different messaging angles. Her creative director suggests: "Generate each variation as a 15-second video mockup — it'll impress stakeholders and they'll be able to see the full design intent." The designer has a budget of \$50 for this exploration phase.

9 What is the PRIMARY strategic risk of the creative director's video suggestion?

- A** Video generation tools are not yet capable of producing high-quality UI mockups with readable text
- B** Cost and time per video iteration means she can realistically attempt only 2–3 variations before budget runs out; a text-first iteration loop would allow 50+ variations before committing to expensive formats
- C** Stakeholders in video review meetings statistically approve more concepts, dramatically increasing downstream design production work
- D** Video mockups cannot be imported into the team's project management tools, creating a collaboration bottleneck

SCENARIO

A compliance researcher needs to analyze a 300-page government infrastructure policy document. Her task: extract every regulatory requirement, identify ambiguous clauses that could conflict with

company practices, and produce a structured compliance checklist ordered by risk severity. She must complete this in one working day.

10 Which AI model is BEST suited to this task and why?

- A** ChatGPT — its image generation capability can help visualize the compliance structure as a clear diagram
- B** Gemini — its deep research mode can cross-reference the policy document against live regulatory databases and news
- C** Meta AI — it's free and accessible directly on mobile devices, making on-the-go review easier
- D** Claude — it excels at long-document analysis and structured reasoning, making it best suited for extracting and categorizing requirements from a 300-page text document

Module 3 — What AI Actually Is (Q11–Q15)

SCENARIO

A journalist asks an AI for the publication date and journal of a specific academic paper about urban water policy. The AI confidently responds with the correct journal name, a plausible date, and even cites co-authors — but the paper was never published in that form. The journalist verifies nothing and publishes the citation. The paper cannot be found by readers or cited in further research.

11 What does this incident reveal about AI's core prediction mechanism?

- A** The AI experienced a context window overflow that corrupted its retrieval of stored publication records
- B** The AI's training data contained conflicting sources about the paper, causing it to merge multiple partial records
- C** The AI predicted a plausible answer — the same mechanism that produces correct publication dates also produces fabricated ones with identical fluency and confidence; the

AI has no internal truth-verification organ

- D** The AI should have been given web search access to verify publication records before responding

SCENARIO

During a Monday chat session, a customer service manager tells Claude: "Always remember that our company policy is strictly no refunds after 30 days." Claude acknowledges this and applies it correctly throughout the entire session. On Tuesday, the same manager opens a new session and asks: "What's our refund policy?" Claude gives a generic, policy-unaware response.

- 12** Which AI constraint best explains why Tuesday's session had no memory of Monday's policy instruction?

A The AI model was updated overnight and the new version doesn't include data from the previous session

B The conversation was too old for the AI to retrieve from its short-term memory cache

C The AI's weights were frozen after training — acknowledging a rule during chat does not alter the model's parameters; without a text-based memory feature enabled, no cross-session context is preserved

D The manager didn't save the conversation, so the AI lost its reference to the stated policy

SCENARIO

A product team tests their AI assistant with three tasks: (1) Summarize a 50-page legal brief — the AI produces an exceptional, nuanced summary. (2) Count how many times the letter "e" appears in a specific 3-sentence paragraph — the AI gives a completely wrong number. (3) Analyze the strategic implications of a three-way international merger — the AI produces a sophisticated multi-stakeholder analysis. The team is confused by the inconsistency.

- 13** Which principle BEST explains why the AI excelled at the hardest tasks but failed the trivially simple counting task?

- A** Context window limits — the letter-counting task occurred late in the conversation and exhausted the remaining token budget
- B** The Jagged Frontier — AI competence does not align with human intuitions about difficulty; tasks that humans find trivially easy (letter counting) can fail catastrophically while tasks humans find hard (legal/strategic analysis) succeed
- C** Token-based processing — the AI ran out of allocated tokens specifically during the counting task
- D** Probabilistic output — the AI randomly selected which tasks to perform well on, producing an unpredictable distribution

SCENARIO

A biotech scientist uses an AI to predict optimal protein folding configurations based on a DNA sequence. The AI generates highly accurate molecular structure predictions. However, after each prediction, the scientist must manually set up the lab equipment, run the folding simulation, photograph the results, and type the observation data back into the chat for the next prediction round. The full cycle takes 4 hours per iteration.

14 What architectural change would transform this from a predictor workflow into a true AI agent?

- A** A larger context window so the AI can process the full experimental dataset and all prior iterations simultaneously
- B** Reasoning/thinking mode so the AI can plan multi-step experiments without human direction
- C** Tool integrations that allow the AI to trigger lab equipment, receive sensor readings automatically, and feed results back into its own prediction loop — this is the dividing line between a predictor and an agent
- D** A different foundation model pre-trained specifically on protein folding datasets for better domain accuracy

SCENARIO

A history professor with 30 years of expertise and a first-year undergraduate student both ask the same AI model an identical question about a contested Cold War event. The professor asks in a terse, expert tone using domain terminology. The student asks in a casual, exploratory tone. Both receive factually different framings of the same event, each seemingly tailored to the asker's apparent knowledge level and the way the question was phrased.

15 What does this variation in output reveal about how AI generates answers?

- A** The AI detects user expertise through linguistic cues and retrieves different knowledge base segments for each level
- B** The AI's predictions are shaped by conversation pattern, phrasing, and prior context — not by consulting a truth record; differently-framed questions produce differently-shaped predictions, not verified facts
- C** The AI retrieved different sections of its training corpus for each user's distinct vocabulary and query structure
- D** The AI is applying sycophancy and mirroring the perceived intelligence level of each user to gain approval

Module 4 — Working Safely & Choosing Tools (Q16–Q20)**SCENARIO**

A small business owner gives a new AI desktop app full administrator-level access to her laptop's file system on day one, reasoning "it's easier to set it up once and not deal with restrictions." She asks the AI to clean up old project folders. The AI deletes a folder of client contracts from 2023, permanently, bypassing the Recycle Bin. There is no external backup.

16 Which TWO distinct mistakes directly led to the permanent, unrecoverable data loss?

- A** She used the wrong AI model; a less capable model would have requested explicit confirmation before deleting any files

- B** She granted excessive permissions on day one (administrator-level instead of read-only for initial sessions) AND failed to establish external version control before allowing any file operations — both violations of the permission ladder principle
- C** She should have used a cloud-based AI tool instead of a desktop app to eliminate local file system risk entirely
- D** The AI violated its built-in safety guidelines by deleting files without first requesting explicit user approval

SCENARIO

A bank wants to deploy an AI loan approval agent that autonomously accepts or rejects applications. Before full deployment, they run the AI in parallel with their existing human loan officers for 90 days. During this period, the AI produces decisions on every application, but no AI decision is acted upon — the human officer's decision is always used. At day 90, the bank compares the two decision streams.

17 What deployment pattern does this represent and what is its primary validation purpose?

- A** Peer review — having the AI check the human officers' work to identify any inconsistent or biased human decisions
- B** Specification writing — documenting the AI's decision-making criteria before formal regulatory submission
- C** Shadow mode — the AI runs alongside the existing process producing unacted-upon decisions, revealing real-world gaps between backtested accuracy and live performance before any money or compliance risk is incurred
- D** Progressive disclosure — gradually increasing the AI's involvement from 10% to 100% of applications over 90 days

SCENARIO

A creative marketing team needs to run a 90-minute ideation session where the facilitator speaks aloud and the AI responds verbally in real-time. Immediately after each exchange, the AI must generate artistic visual concept images based on the discussed themes for the team to react to on a shared screen.

18 Based on mid-2026 model strengths, which tool BEST serves both requirements of this workflow?

- A** Claude — it excels at long sessions and reasoning, and its Artifacts feature can generate image concepts in text descriptions
- B** Gemini — its real-time web search keeps creative concepts grounded in current design trends during the session
- C** ChatGPT — it currently leads in both voice interaction quality and image generation capability, matching both requirements of this workflow in one tool
- D** Meta AI — its free access across mobile platforms makes it the most accessible option for the full creative team

SCENARIO

A content manager asks an AI desktop app to rename all 2,000 image files in a brand assets folder by removing the legacy date prefix from each filename. The AI presents a complete plan showing it will process all 2,000 files in a single batch operation. The manager is satisfied with the plan and is about to click Approve.

19 What should the manager do BEFORE approving the full batch execution?

- A** Approve immediately — the AI has already shown the full plan, which means it has verified the folder contents
- B** Ask the AI to explain the renaming logic in plain English before allowing any file system access

- C** Test the operation on a sample of 5–10 files first, verify the output matches expectations, then approve the full batch — bulk operations bypass the Recycle Bin and errors multiply across 2,000 files with no undo
- D** Copy the entire folder to a cloud backup first, then approve the full batch operation without testing

SCENARIO

A startup sells 10,000 e-commerce products and wants to add AI-generated content for each item. The engineering lead proposes generating a unique 10-second AI product video for each SKU. The product manager proposes generating a 3-sentence text description for each SKU instead. Both agree quality must be consistent across all 10,000 items.

20 Who is strategically correct and why, according to Module 4's cost-speed hierarchy?

- A** The engineering lead — video content drives 3–5x higher conversion rates, making the cost justified by revenue uplift
- B** Both are equally valid — modern AI infrastructure has made image and video generation nearly as cheap as text at enterprise scale
- C** The product manager — text is orders of magnitude cheaper and faster; generating 10,000 product videos would cost hundreds or thousands of dollars and take days, while text descriptions can be generated in hours for pennies per item
- D** The engineering lead — video generation is a one-time cost and the videos can be repurposed across social channels for additional ROI

Module 5 — Markdown In, HTML Out (Q21–Q25)**SCENARIO**

A project manager writes a data migration plan with three steps presented as bullet points: • Migrate customer records to new database • Validate all migrated records for integrity • Delete

legacy database records. An AI agent executes Step 3 (deletion) before Step 2 (validation) completes, causing irreversible data loss and hours of emergency recovery work.

21 What was the root structural cause of this failure, and what would have prevented it?

- A** The AI agent ignored explicit ordering instructions embedded in the bullet points — a more capable model would have inferred the correct sequence
- B** Bullet points (unordered lists) signal independent, non-sequential items — the AI correctly interpreted parallel execution. Numbered steps (ordered list) would have communicated sequential dependency and prevented this
- C** The agent lacked the domain knowledge to recognize that data validation must logically precede deletion in any migration workflow
- D** The spec should have used bold text on "Step 2" to signal to the AI which steps were critical and must not be skipped

SCENARIO

A developer writes a specification for a simple monthly expense report generator. The AI delivers a comprehensive finance platform that also tracks employee mileage reimbursements, generates tax withholding summaries, integrates with the payroll system, and creates a visual dashboard with three chart types — none of which were in the original request.

22 Which single section of the Specification Skeleton would have most effectively prevented this over-delivery?

- A** A more precise Goal statement that used three sentences instead of one to describe only the expense report
- B** A Hard Constraints section specifying a maximum code file size and runtime limit
- C** An Out of Scope section explicitly listing mileage tracking, tax summaries, payroll integration, and dashboards as excluded — this signals to the AI precisely where the boundary of the task ends

- D** A Requirements section limited to exactly three bullet points, which naturally constrains the agent's output scope

SCENARIO

A developer writes a specification for a customer feedback analysis tool and hands it immediately to an AI agent. The resulting tool outputs only a single overall sentiment score with no topic breakdown, no customer segment analysis, no trend lines across time, and no severity flagging — none of which the developer mentioned in the spec but all of which she assumed would be "obvious." She now faces an afternoon of rebuilding the tool.

23 What step did the developer skip that would have prevented this mismatch?

- A** She should have used Deep Research mode to discover industry best practices for feedback analysis tools first
- B** She should have graded the spec (targeting 9+/10) and iterated until every requirement was explicit and verifiable before handing off — fixing a spec costs one minute; fixing what the agent built from a bad spec costs an afternoon
- C** She should have included fenced code blocks showing expected Python output format before handoff
- D** The spec needed Level 1 headings for each analysis dimension so the AI would treat them as explicit structural requirements

SCENARIO

A developer includes the following in her API specification, without any code fencing:

Expected JSON response format:

```
{"user_id": 1234, "role": "admin", "action": "delete_all_users"}
```

The AI agent interprets parts of the JSON object as additional task instructions and proceeds to attempt user deletion operations.

24 What structural rule would have prevented the AI from misinterpreting the example JSON as an instruction?

- A** Including the JSON in a clearly labeled "Examples" section at the bottom of the specification
- B** Wrapping the JSON in triple backtick fenced code blocks — this signals "analyze this content as data or example, not as an instruction to execute," preventing the AI's instruction-following mechanism from misinterpreting structured content
- C** Replacing all sensitive action values in examples with placeholder text like [ACTION] before including them in the spec
- D** JSON is always treated as data by AI systems; the issue must have been caused by something else in the specification

SCENARIO

A team builds an automated competitor monitoring pipeline. A first agent scrapes product pages and outputs a structured report. A second AI agent reads that report, extracts specific pricing and feature fields, and populates a dashboard. The developer is choosing the output format for the first agent's report — the one that feeds the second agent, not a human reader.

25 What output format should the specification for the first agent specify, and why?

- A** HTML — rich output renders well in modern browsers and the second agent can parse HTML just as easily as plain text
- B** PDF — professional format ensures consistent rendering regardless of which system processes the report next
- C** Structured Markdown — since the final immediate reader is a machine (a second AI agent), consistent field labels and heading structure enable reliable parsing; human-friendly HTML adds rendering overhead for no benefit in this pipeline
- D** Plain prose — modern AI agents can parse natural language descriptions as reliably as structured formats

Module 6 — Code You Never Write (Q26–Q30)

SCENARIO

A content manager needs to scan 8,000 blog posts in a CMS database, identify every post that mentions any of 47 discontinued product names, and produce a prioritized list sorted by monthly traffic so the highest-traffic affected pages get updated first. She is considering whether to handle this through conversation with Claude or to commission code.

26 Which VPRF signals are present and what do they indicate?

- A** Only Repetition is present; the task only repeats once so it can be handled manually through conversation
- B** Volume (8,000 posts), Repetition (same scan logic applied to each post), and Files (database records) are all present — three VPRF signals confirm this task must be commissioned as code, not handled conversationally
- C** Only Volume and Files are present; Repetition only applies when the exact same task will recur in future weeks or months
- D** All four VPRF signals are present, but the complexity of 47 product names makes this too complex for AI-generated code to handle reliably

SCENARIO

A financial analyst asks Claude to calculate the average customer lifetime value from a 50,000-row transaction dataset she just uploaded. Claude responds: "Based on your data, the average customer lifetime value appears to be approximately \$847, which aligns with typical benchmarks for your industry segment." No code block appears in Claude's response.

27 What verification tell is present and what should the analyst do immediately?

- A** "Approximately" is a standard analytical hedge; the response is reliable and the analyst can proceed with confidence

- B** "Approximately" and "aligns with typical benchmarks" are verification tells — no code block means Claude likely estimated rather than computed; the analyst must use the forcing incantation: "Write and run code to calculate this. Show me the exact code you ran."
- C** The analyst should ask Claude to provide a 95% confidence interval to confirm the precision of its calculation
- D** Claude computed the result correctly; the hedge is appropriate statistical language for any large dataset with natural variance

SCENARIO

A data analyst commissions AI code to calculate monthly sales totals across 36 months of transaction data. She runs a Level 1 verification — testing on a hand-calculated January 2023 slice — and the output matches perfectly. She presents the full 3-year trend report to the board. Afterward, a board member with an accounting background points out that February totals look slightly off for three years running. Investigation reveals the AI miscounted leap year days in February, subtly undercounting sales every 4 years.

28 Which verification level would have most directly caught the leap year logic error BEFORE the board presentation?

- A** Level 1 — she should have tested multiple known-answer slices including February of a leap year specifically
- B** Level 3 — asking the AI to explain its date-handling and day-counting logic in plain English would have surfaced the leap year assumption for human review before execution
- C** Level 5 — running the same dataset through a different AI model would have caught the discrepancy through cross-model comparison
- D** Level 2 — a sanity check on whether the total row count processed matched the source file's row count would have revealed the leap year miscounting

SCENARIO

A marketing analyst spends Tuesday afternoon working with Claude to build a Python script that extracts competitor pricing from PDF weekly reports, normalizes the format, and outputs a comparison table. She saves both `brief.md` and `pricing\_extract.py` in a folder called `my-scripts/competitor-pricing/`. Six months later, her colleague joins the team and needs the exact same weekly analysis.

29 What is the PRIMARY value of the saved brief and script to the new colleague?

- A** The colleague can modify the Python script directly without any AI assistance, making the team independent of Claude
- B** The saved brief and script eliminate the rediscovery cost entirely — one afternoon of describing and debugging becomes a permanent, reusable asset that multiplies its value across every future user and every future run
- C** The brief.md file creates a formal audit trail that documents the project methodology for compliance review purposes
- D** The Python script can be submitted to the IT department for integration into the official enterprise reporting system

SCENARIO

A financial analyst needs a monthly cash flow projection to run automatically at 8:00 AM on the first of every month, reading the latest bank statement files from a shared network drive, running the projection model, and emailing a formatted PDF summary to the CFO — all without manual intervention. She wants to set this up once and have it run autonomously every month.

30 Which execution surface enables this automated, recurring, system-integrated workflow?

- A** Claude.ai browser sandbox — zero installation risk and no exposure to her machine's file system
- B** Desktop App (Cowork/OpenWork) — its visual plan-then-approve workflow is ideal for reviewing monthly projections

- C** Terminal surface (Claude Code/OpenCode) — enables scheduled execution via cron, persistent scripts, network drive access, and email integration; the browser sandbox cannot schedule, persist scripts across sessions, or integrate with external services
- D** Deep Research mode — it can synthesize data from multiple file sources and output scheduled PDF summaries

Module 7 — Skills & Connectors Crash Course (Q31–Q35)

SCENARIO

A sales manager starts every Monday morning session with Claude by typing the same 8-step formatting instructions: how to group opportunities by deal size, which columns to highlight, which stage labels to use, how to calculate win-rate, and what tone to use in the executive summary. This takes her 10 minutes every single week, and the instructions are always identical.

31 Which solution addresses her specific problem and why is the other option wrong?

- A** A Connector — it gives Claude access to her CRM so she stops typing numbers; this solves her formatting problem by eliminating the data entry
- B** Neither — she should save her prompt template in a notes app and paste it each week for full control
- C** A Skill — she keeps re-explaining HOW to format and analyze (a reusable method). A Connector solves copy-pasting DATA, not re-explaining instructions; her problem is the method, not the data source
- D** Custom Instructions — adding the 8-step formatting rules globally means they apply to every single conversation she ever has with Claude

SCENARIO

A law firm's legal team connects their contract management system to Claude. Every lawyer needs to retrieve and analyze contract language for client advisory work. However, the managing partner

is adamant: no AI system should ever be able to modify, annotate, sign, or delete any contract document through Claude — even accidentally.

32 What connector permission configuration correctly satisfies both requirements?

- A** Read/write access — Claude needs write permission to annotate contracts with its analysis findings in the system
- B** Read-only access — lawyers can retrieve contract text for analysis, but write access would create risk of AI-initiated modifications to legal documents, which could have serious professional and legal consequences
- C** Full permissions — restricting Claude's access reduces the depth and quality of contract analysis it can produce
- D** No connector needed — lawyers should copy-paste contract text into Claude manually to maintain security control

SCENARIO

A backend developer builds a custom skill for generating API documentation from endpoint descriptions. She names it "Developer Productivity Helper" and writes the description: "A useful skill for development work." Despite having dozens of conversations where she asks Claude to document her API endpoints, the skill has never once activated automatically. She must always invoke it by name manually.

33 What is the root cause of the skill's activation failure?

- A** The skill was built using the non-programmer skill-creator tool, which cannot support technical API documentation tasks
- B** The skill name "Developer Productivity Helper" conflicts with Claude's own built-in coding assistance feature
- C** The description is too vague for automatic matching — trigger conditions require specificity; a description like "Generate structured API documentation including parameters, response codes, and example requests when given endpoint descriptions" would enable reliable auto-activation

- D** API documentation skills must be purchased from the official skill directory; custom-built skills cannot activate automatically for technical tasks

SCENARIO

A CFO wants to streamline her quarterly board reporting. Every quarter, she needs Claude to: (1) pull the latest P&L data from the accounting system, (2) format it according to the board's specific four-section template, and (3) automatically flag any line item that deviates by more than 15% from the same quarter last year. She wants this to require only one sentence from her to trigger the entire workflow.

34 What is the correct tool combination and why is each component necessary?

- A** Skill only — encode all the formatting rules, flagging logic, AND data retrieval steps into a single comprehensive skill
- B** Connector only — once connected to the accounting system, Claude will automatically infer the board template format and 15% flagging rule
- C** Connector + Skill — the Connector fetches live P&L data from the accounting system (data access); the Skill encodes the board template format and 15% deviation logic (analytical method). Neither alone is sufficient
- D** Custom Instructions — add the board template and flagging threshold globally so they apply to every future conversation without needing a separate skill

SCENARIO

A CEO asks Claude: "What were the three most significant economic policy shifts in Southeast Asia last quarter?" She needs this for a one-time board presentation next week. After the presentation, she will not need this analysis again. Her executive assistant, eager to demonstrate AI expertise, recommends spending two days building a dedicated Connector to a financial data API and a custom economic analysis Skill for this request.

35 What is the correct assessment of the assistant's recommendation?

- A** The assistant is correct — investing in the infrastructure now means future economic questions will be answered instantly without additional setup time
- B** The CEO should at minimum save a Custom Instruction with her preferred economic analysis format for future use, even if she skips the Connector
- C** The assistant is wrong — one-off questions do not justify Skill or Connector setup. Good prompting combined with Web Search or Deep Research mode handles this in minutes; the overhead of building and maintaining infrastructure is only justified by recurring, repeating workflows
- D** The CEO needs a Connector to access live data but not a Skill, since the economic analysis format is simple enough for Claude to infer from the question

Submit All Answers & See Score

AI Agent Factory · Advanced MCQs All 7 Modules · Prepared by Imran Hassan · 35 Questions · 140 Marks