

The Agent Factory Thesis — MCQs

Panaversity Agent Factory AI Prompting 2026

Prepared by: Imran Hassan

0

SCORE

0

ANSWERED

50

TOTAL

Part 1 — Core Thesis & Paradigm Shift (Q1–Q10)

SCENARIO

A venture capitalist is evaluating two companies: Company A sells CRM software on a per-seat subscription. Company B deploys AI Workers that handle lead generation, outreach, and CRM hygiene as outcomes. According to the Agent Factory thesis, which model represents the future?

1 What is the core claim of the Agent Factory thesis about future valuable companies?

A They will build better SaaS software with AI features

B They will manufacture AI employees rather than sell traditional software

C They will replace all human employees with robots

D They will sell AI models on a per-token pricing model

SCENARIO

A startup founder says: "We sell our platform to enterprise clients who pay \$500/user/month." His advisor, familiar with the Agent Factory thesis, says this model is outdated. What commercial model does the thesis describe?

2 How does the Agent Factory thesis redefine the commercial relationship between AI-Native companies and their clients?

- A "You don't buy from these companies. You hire them."
- B "You don't hire from these companies. You buy from them."
- C "You subscribe to these companies on a monthly basis."
- D "You partner with these companies for co-development."

SCENARIO

A traditional SaaS company charges per seat and requires humans to manually operate the software. An AI-Native company charges based on results delivered. A student is mapping these differences to the Agent Factory thesis paradigm shift table.

3 In the SaaS to Agent Factory paradigm shift, what changes in the "Value Metric"?

- A From per-outcome results to per-seat subscriptions
- B From per-seat subscriptions to per-outcome results
- C From annual contracts to monthly billing
- D From token pricing to feature pricing

SCENARIO

In the old SaaS era, humans manually operated tools. A manager asks: "In the Agent Factory era, what is my new role?" According to the thesis, the human role fundamentally changes.

4 How does the human role change from the SaaS era to the Agent Factory era?

A From Supervisor to Operator

B From Operator to Supervisor & Verifier

C From Manager to Individual Contributor

D From Developer to Designer

SCENARIO

A product manager reads the thesis and sees the term "Digital FTEs." She's confused — are these bots, scripts, or something else? She wants to understand exactly what Digital FTEs are.

5 What are "Digital FTEs" or "AI Workers" in the Agent Factory thesis?

A Simple chatbots that answer FAQs

B Role-based agents that are hired, assigned, rostered, and retired within AI-Native companies

C Human contractors working remotely using AI tools

D Software scripts that run on a schedule

SCENARIO

Cursor, a popular coding tool, published data in February 2026 that the Agent Factory thesis cites as evidence for the 10-80-10 rule in action. A student is looking for this specific statistic.

6 What evidence did Cursor report in February 2026 supporting the Agent Factory thesis?

A 50% of developers prefer AI coding over manual coding

- B** 35% of merged pull requests originated from autonomous cloud-based agents
- C** AI reduced software development time by 80%
- D** 100% of bug fixes were done by AI in Q1 2026

SCENARIO

A CEO wants to implement the 10-80-10 operating rhythm in her company. She needs to understand what each percentage represents before restructuring her team's workflow.

7 In the 10-80-10 operating rhythm, what does each segment represent?

- A** 10% planning, 80% coding, 10% testing
- B** First 10% Human defines spec → Middle 80% AI executes autonomously → Final 10% Human reviews and approves
- C** 10% AI involvement, 80% human work, 10% quality check
- D** 10% budget allocation, 80% execution, 10% profit

SCENARIO

The thesis describes how AI Workers will eventually go beyond just executing tasks — they will autonomously purchase compute, negotiate API contracts, and provision cloud services. A student needs to understand what this represents.

8 What does "Agents as Economic Actors" mean in the Agent Factory thesis?

- A** Agents will replace economists in financial institutions
- B** AI Workers transition from tool execution to autonomous market participation — buying compute, negotiating contracts, provisioning services within budget envelopes

C Agents will manage stock portfolios for clients

D AI will replace human economists in policy-making

SCENARIO

A data center company is seeing unprecedented demand. An investor asks: "What's driving this massive infrastructure buildout?" The thesis cites specific 2025-2026 infrastructure numbers to explain the AI boom.

9 What does the 2025-2026 infrastructure data in the thesis indicate about the AI era?

A Office construction is booming as companies expand their workforce

B US data center construction (\$42B/year) now nearly equals office construction (\$44B), with Big Tech spending \$600B+ on AI infrastructure

C AI infrastructure spending has plateaued at \$100 billion

D Residential construction is shifting toward smart home AI integration

SCENARIO

A worker is worried: "Will AI take my job completely?" The thesis provides a nuanced answer about what happens to professional roles in the AI era, quoting a specific phrase about developers.

10 According to the Agent Factory thesis, what happens to the developer's role in the AI era?

A "The developer disappears completely"

B "The developer doesn't disappear. The developer does more."

C "The developer becomes a manager of AI systems only"

D "The developer's role reduces by 80%"

Part 2 — Seven Principles of Problem Solving (Q11–Q20)

SCENARIO

An AI agent is asked to "help with the project." It responds with a long description of what could be done, but takes no actual action. According to the Seven Principles, this agent is violating the first principle.

11 What does "Bash is the Key" mean as the first principle of General Agent Problem Solving?

A Agents must always use Bash scripting language

B Agent capacity for action, not just description — agents must execute, not just explain

C Agents should respond quickly like a bash (hit)

D Bash terminal must be the primary interface for all agents

SCENARIO

An agent is working on a large complex task. Instead of breaking it into small steps, it attempts the entire task in one massive operation. If it fails, nothing can be undone. This violates one of the Seven Principles.

12 Which principle requires agents to work in small, reversible steps that enable undo capability?

A Bash is the Key

B Code as Universal Interface

C Small, Reversible Decomposition

D Observability

SCENARIO

An agent stores all its work progress in the conversation context. When the conversation ends, all the work is lost and must be redone from scratch. This violates one of the Seven Principles.

13 Which principle addresses the difference between volatile conversation and durable filesystem storage?

A Verification as Core Step

B Persisting State in Files

C Constraints and Safety

D Observability

SCENARIO

An agent performs several actions in the background without any visibility into what it's doing. The human supervisor has no way to audit or understand what happened. This violates the seventh principle.

14 What does the Observability principle require in agent problem-solving?

A Agents must observe users to personalize responses

B Complete visibility into agent actions — no hidden processes allowed

C Agents should use observation cameras in physical spaces

D Agents must observe competitors' strategies

SCENARIO

An agent produces an output and immediately sends it to production without checking it. Later, it's discovered the output was incorrect and caused system errors. This violates one of the Seven Principles.

15 Which principle states that all meaningful outputs must be checked before deployment?

A Bash is the Key

B Verification as Core Step

C Persisting State in Files

D Small Reversible Decomposition

SCENARIO

An agent starts operating with maximum autonomy from day one — making decisions, deleting files, sending emails — without any explicit human-defined permissions. This violates one of the Seven Principles.

16 What does the "Constraints and Safety" principle state about autonomy?

A Agents should have unlimited autonomy to maximize efficiency

B Explicit permissions must be set; autonomy is earned per task, not granted upfront

C Agents should request permission for every single micro-action

- D** Autonomy is determined by the AI model's capability, not by humans

SCENARIO

A team uses Mode 1 (Problem-Solving engagement) for their daily tasks. Domain experts on the team are deciding which tool to use — Claude Code or Claude Cowork. The thesis provides guidance on this.

- 17** In Mode 1 (Problem-Solving), which tools are used by engineers vs. domain experts?

- A** Engineers use Cowork; domain experts use Code variants
- B** Engineers use Code variants (Claude Code/OpenCode); domain experts use Cowork/OpenWork
- C** Both engineers and domain experts use the same tools
- D** Engineers use ChatGPT; domain experts use Gemini

SCENARIO

A team member asks: "What's the key difference between Mode 1 and Mode 2 agent engagement?" He wants to understand when to use each mode and what persists after the session ends.

- 18** What is the KEY difference between Mode 1 (Problem-Solving) and Mode 2 (Manufacturing) engagement?

- A** Mode 1 uses AI; Mode 2 uses human workers
- B** Mode 1 delivers outcome within a session (nothing persists); Mode 2 produces a persistent AI Worker that survives the session and joins the workforce
- C** Mode 1 is for simple tasks; Mode 2 is for complex tasks

D Mode 1 is free; Mode 2 requires enterprise subscription

SCENARIO

A developer needs to communicate precise technical requirements to an AI agent. Instead of natural language descriptions, the thesis recommends a specific approach for maximum precision.

19 What does "Code as Universal Interface" principle recommend for communicating with AI agents?

A Always use natural language for human-AI communication

B Use schemas, tables, and code blocks for precision in communication

C Use diagrams and flowcharts as the primary interface

D Use voice commands for faster agent interaction

SCENARIO

In Mode 2 (Manufacturing), a team always uses the same tool regardless of which team member is operating it. The thesis specifies which tool must always be used for manufacturing engagement.

20 Which tool is ALWAYS used in Mode 2 (Manufacturing engagement), regardless of the operator?

A Claude Cowork or OpenWork

B Claude Code or OpenCode

C ChatGPT or Gemini

D Any tool the operator prefers

Part 3 — Seven Invariants of Agent Factory (Q21–Q36)

SCENARIO

A company deploys AI Workers that operate completely autonomously — no human defines the intent, no one sets the budget or authority envelope. The thesis warns this is dangerous. Why?

21 What does Invariant 1 ("Human is Principal") require, and what happens if it's violated?

- A** Humans must write all the code; violation causes slow development
- B** Every action chain must originate with a human defining intent, budget, and authority; violation causes unowned systems and unaccountable outcomes
- C** Humans must approve every micro-action; violation causes agent inefficiency
- D** The Principal must be a C-level executive; violation causes governance issues

SCENARIO

A company has deployed an AI workforce but has no personal delegate agent for each human supervisor. The humans must manually orchestrate every worker interaction. The thesis identifies this as a specific failure mode.

22 What does Invariant 2 ("Every Human Needs a Delegate") solve, and what's the current tool realizing it?

- A** It solves budget management; current tool is Paperclip
- B** It solves the human bottleneck problem; a personal agent (current: OpenClaw) holds context, carries authority, and brokers downstream work
- C** It solves worker hiring; current tool is Inngest

D It solves data storage; current tool is Postgres

SCENARIO

An AI workforce operates without a management layer. Workers collide on the same tasks, budgets leak without enforcement, and there's no audit trail. The thesis identifies this as a specific invariant violation.

23 What is Paperclip's role in the Agent Factory architecture (Invariant 3)?

A It serves as the event nervous system for agent coordination

B It exposes the workforce lifecycle — hire, assign, govern, observe, retire — as callable APIs

C It manages database connections for AI Workers

D It provides the personal delegate for each human supervisor

SCENARIO

A company forces all its AI Workers to run on the same engine — Claude Managed Agents — regardless of whether the task needs maximum reliability or just routine work. The thesis warns this approach causes a specific problem.

24 What does Invariant 4 ("Each Worker Picks Own Engine") prevent, and what engines are available?

A It prevents data loss; engines include SQL databases and file systems

B It prevents uniform trade-offs; engines include Dapr Agents, Claude Managed Agents, OpenAI Agents SDK, Cursor SDK

C It prevents budget overruns; engines include cost calculators

- D** It prevents agent collisions; engines include traffic controllers

SCENARIO

An AI Worker is generating financial reports but instead of reading from the company's actual database, it "invents" plausible-sounding numbers from its training. The audit trail is gone and liability is untraceable. Which invariant was violated?

25 What does Invariant 5 ("Worker Runs Against System of Record") prevent?

- A** It prevents workers from running too fast
- B** It prevents hallucinations and untraceable liability by requiring workers to read/write from authoritative databases rather than inventing facts
- C** It prevents duplicate records in the database
- D** It prevents workers from accessing sensitive data

SCENARIO

The thesis recommends Postgres as the "center of gravity" for operational storage in AI-Native companies. A developer asks why not use specialized databases for everything.

26 Why does the Agent Factory thesis recommend Postgres as the default System of Record?

- A** It is the cheapest database option available
- B** Postgres handles documents (JSONB), search (full-text), vectors (pgvector), work queues, and analytics — consolidate by default, specialize deliberately
- C** Anthropic specifically built Claude to work only with Postgres

D Postgres is the only open-source database compatible with MCP

SCENARIO

An AI workforce has a fixed roster of 10 workers. When business demand spikes and 50 workers are needed for a specific project, humans must manually create and configure each new worker — causing massive delays. Which invariant is being violated?

27 What does Invariant 6 ("Workforce Expandable Under Policy") enable?

A Workers can expand their own skills over time

B Authorized agents can automatically hire new workers within the authority envelope without human intervention — enabling scale

C The workforce budget automatically expands as needed

D Workers can expand their memory storage capacity automatically

SCENARIO

The thesis provides a critical reliability calculation: a 10-step task at 95% per-step reliability. Without memoization, this gives 74% completion. With memoization, it reaches 99.7%. This is related to Invariant 7.

28 What does Invariant 7 (Nervous System) address, and what is the current tool realizing it?

A It addresses agent intelligence; current tool is Claude Opus

B It addresses events, durability, and flow — external triggers, worker-to-worker handoffs, step memoization, retry; current tool is Inngest

C It addresses UI design; current tool is Figma

D It addresses security; current tool is Auth0

SCENARIO

Without Invariant 7's nervous system, what happens to agent coordination? The thesis explains three specific failure modes when this invariant is absent.

29 What are the failure consequences when Invariant 7 (Nervous System) is absent?

A Agents become too expensive and budget overruns occur

B System runs at typing speed; workers can't coordinate without human routing; multi-step runs lose work to crashes; one spike starves all others

C Agents lose their memory between conversations

D Workers refuse to execute tasks without human confirmation

SCENARIO

The thesis describes the two-layer model — Edge Layer and AI Workforce Layer. A student tries to understand why you can't just have one layer without the other.

30 What happens if the Edge Layer and AI Workforce Layer operate independently without each other?

A Both layers work fine independently — they are optional enhancements

B Personal agents without workforce = constrained assistants; workforce without edge agents = forces manual human orchestration at typing speed

C The Edge Layer can operate alone but the Workforce Layer needs both

- D** Both layers can function independently but produce lower quality outputs

SCENARIO

A team is choosing between Dapr Agents and OpenAI Agents SDK for a critical financial processing worker that absolutely cannot fail. The thesis provides a clear engine selection framework.

- 31** According to the engine selection framework, which engine should be used for a task that "can't fail"?

- A** Claude Managed Agents — fully managed and reliable
- B** Dapr Agents — with durable execution for zero failure tolerance
- C** Cursor SDK — proven at parallel cloud scale
- D** OpenAI Agents SDK — highest model quality

SCENARIO

A company wants minimal operational burden — they don't want to manage infrastructure but also need high reliability. According to the engine selection framework, what's the right choice?

- 32** Which engine is best for "shouldn't fail, minimal ops" according to the Agent Factory framework?

- A** Dapr Agents
- B** Claude Managed Agents — fully managed runtime
- C** OpenAI Agents SDK
- D** Cursor SDK

SCENARIO

A company is worried about vendor lock-in but still needs a reliable agent engine. They want the ability to switch AI providers in the future without rebuilding everything.

33 Which engine has zero vendor lock-in (Apache 2.0 license) according to the thesis?

- A** Claude Managed Agents
- B** OpenAI Agents SDK
- C** Dapr Agents — zero lock-in with Apache 2.0
- D** Cursor SDK

SCENARIO

An AI-Native company needs to expose its System of Record (CRM, ERP, ticketing) to AI Workers while maintaining policy governance. The thesis identifies the universal standard enabling this.

34 What standard does MCP (Model Context Protocol) enable in the Agent Factory architecture?

- A** Point-to-point API connections between specific tools
- B** Universal standard enabling Worker-to-tool communication and exposing Systems of Record under policy governance
- C** A payment protocol for AI economic transactions
- D** A security standard for encrypting agent communications

SCENARIO

The thesis quotes a key phrase encapsulating the entire three-step operational model: human defines intent, agents execute, humans verify. A student needs to identify this exact quote.

35 What is the three-step operational model quoted in the Agent Factory thesis?

A "Plan, Execute, Review"

B "Humans define intent. Agents execute. Humans verify outcomes."

C "AI proposes. Humans approve. Systems deploy."

D "Design. Build. Ship."

SCENARIO

The thesis makes a distinction between what stays stable (invariants) and what will change over time (implementations). A student needs to identify which elements are truly stable vs. which will change.

36 According to the thesis, what will **CHANGE** over time vs. what remains **STABLE** (invariant)?

A The Seven Invariants will change; specific tools like Paperclip and Inngest will stay

B Invariants (human principal, management layer, System of Record, etc.) stay stable; tools/products (Paperclip, Inngest, Postgres, etc.) will change

C Everything changes — the thesis is a snapshot with no durable principles

D Specific tools stay stable; the principles evolve annually

Part 4 — Economic Protocols, Workforce & Future (Q37–Q50)

SCENARIO

An AI Worker needs to autonomously purchase compute resources during a task. Four open protocols exist that enable AI agents to make autonomous economic transactions. A student is mapping which protocol belongs to which company.

37 Which economic payment protocol is associated with OpenAI and Stripe?

A AP2

B ACP (Agent Commerce Protocol)

C x402

D MPP

SCENARIO

A blockchain-focused startup is building an AI agent that needs to make crypto-native payments. They want to use a protocol that bridges mainstream and crypto payments. Which protocol should they use?

38 Which economic payment protocol is crypto-native and associated with Coinbase?

A ACP

B AP2

C x402

D MPP

SCENARIO

An AI agent needs to make micropayments — tiny per-use charges for each API call it makes. The thesis identifies a specific protocol designed for consumption-based micropayment commerce.

39 Which protocol handles micropayments for consumption-based commerce, associated with Stripe and Tempo?

A ACP

B AP2

C x402

D MPP (Micropayment Protocol)

SCENARIO

A sales team wants to use AI Workers for their Go-To-Market operations. They want to understand which specific GTM tasks AI Workers can handle according to the thesis.

40 Which of the following is a GTM (Go-To-Market) AI Worker task listed in the Agent Factory thesis?

A Code review and deployment

B Lead enrichment, outreach sequencing, CRM hygiene, pipeline analysis, proposal generation

C HR sourcing and legal redlining

D Financial close and accounts payable processing

SCENARIO

The World Economic Forum has made a projection about job opportunities created by the AI era by

2030. The Agent Factory thesis cites this to argue AI creates opportunities, not just eliminates jobs.

41 What does the World Economic Forum project about job opportunities by 2030 in the AI era?

- A** 500 million jobs will be eliminated with no replacements
- B** 78 million new job opportunities by 2030, requiring 59 of 100 workers to upskill/reskill
- C** Only technical workers will have job security
- D** 25 million jobs will be created exclusively in AI development

SCENARIO

The thesis describes a future trajectory called "Physical AI Workers." A robotics engineer wants to understand how the same Agent Factory architecture extends to physical robots and autonomous vehicles.

42 How does the Agent Factory thesis describe "Physical AI Workers" as a future trajectory?

- A** Physical robots will require completely new management architectures different from digital agents
- B** The same manufacturing architecture extends to embodied agents (robots, autonomous vehicles) running the same authority envelopes, management layers, and runtime engines (with actuators added)
- C** Physical AI Workers will replace the need for digital AI Workers entirely
- D** Physical AI Workers will be governed separately from digital workers with different invariants

SCENARIO

The thesis describes a future where AI Workers might accumulate capital, enter contracts, and work for multiple companies simultaneously. A student wants to understand what this future state is called.

43 What does "Cross-Company Workforce Mobility" mean as a future trajectory in the thesis?

- A** Human employees moving between AI companies
- B** AI Workers gaining portability — hired into one company, transferred to another, serving multiple simultaneously — creating a labor market for AI Workers with rates, reputations, and specializations
- C** AI models being moved between different cloud providers
- D** Software being portable across different operating systems

SCENARIO

The thesis distinguishes between what remains constant (invariants) and what is implementation detail. The quote "The building stands even when the furniture changes" captures this principle.

44 What does "The building stands even when the furniture changes" mean in the Agent Factory context?

- A** Physical office buildings will outlast digital tools
- B** The Seven Invariants (building) remain stable even as specific tools/products (furniture) change over time
- C** Company culture remains stable even as employees change
- D** Cloud infrastructure outlasts the applications running on it

SCENARIO

An AI Worker needs to read from a CRM, collaborate with another AI Worker on analysis, and then write results back. The thesis describes a trust focus area for agents becoming economic actors.

45 As AI Workers become economic actors, what is the primary trust layer focus identified in the thesis?

- A** Model capability improvement and reasoning quality
- B** Mandate enforcement, audit trails, and liability frameworks — not just capability
- C** UI/UX design for better human-agent interaction
- D** Speed optimization for faster agent responses

SCENARIO

A legal team is considering AI Workers for their department. They want to know exactly which legal tasks the Agent Factory thesis identifies as appropriate for AI Workers.

46 Which legal tasks does the Agent Factory thesis identify as suitable for AI Legal Workers?

- A** Court representation and litigation strategy
- B** Document review, redlining, and intake processing
- C** Setting legal precedents and writing legislation
- D** Serving as expert witnesses in trials

SCENARIO

The Agent Factory defines a specific term for the high-level plan that AI Workers receive before starting work — it includes goals, constraints, budgets, and permissions. A student needs to identify

this term.

47 In the Production Engine architecture, what is "Intent"?

- A** The AI model's understanding of natural language
- B** The high-level blueprint encompassing goals, constraints, budgets, and permissions — the human's input to the production engine
- C** The AI Worker's motivation for completing a task
- D** The final outcome delivered by the AI Worker

SCENARIO

The Agent Factory thesis describes how AI "unbundles jobs into tasks." An HR manager is curious what this means for how future work roles will be structured.

48 What does the thesis mean when it says AI "unbundles jobs into tasks"?

- A** AI breaks complex software into smaller microservices
- B** AI separates jobs into individual task components, creating new role combinations — for example, GTM Workers handling specific sales tasks that previously required multiple human roles
- C** AI removes the need for job descriptions in companies
- D** AI makes every employee a generalist by removing specialization

SCENARIO

An AP2 protocol is cited in the thesis as a cross-vendor standard for AI economic transactions. A student wants to know which company is behind this protocol.

49 Which company's economic payment protocol uses cryptographically signed mandates as a cross-vendor standard?

A OpenAI (ACP)

B Google (AP2)

C Coinbase (x402)

D Stripe (MPP)

SCENARIO

A student is reviewing the entire Agent Factory thesis and wants to identify the single most foundational claim — the thesis statement that everything else builds upon.

50 What is the foundational thesis statement of the Agent Factory document?

A "In the AI era, the most powerful companies will build the best AI models"

B "In the AI era, the most valuable companies won't sell software — they'll manufacture AI employees"

C "In the AI era, every company must become a software company"

D "In the AI era, automation will replace 90% of white-collar jobs"