

AI Agent Factory

Foundations Course — agentfactory.panaversity.org

PREPARED FOR
Imran Hassan

MODULE 3 — WHAT AI ACTUALLY IS (CRASH COURSE)

Total Questions: 36 Parts: 3 Concepts Covered: 9 Marks Each: 4 Total Marks: 144

Core Quote to Remember: "It is a prediction machine that learned by reading and has no organ for truth, so it is fluent everywhere, reliable only where the text was thick, and you are the part that checks."

PART 1 THE MACHINE — THREE FOUNDATIONAL IDEAS (CONCEPTS 1-3)

1

IDEA 1 — PREDICTION OVER RETRIEVAL

Imran types a question into AI and expects it to look up the correct answer from a database like Google. But according to this module, what is AI actually doing?

- A Searching through a massive verified database of facts
- B Predicting the most plausible next text — like an advanced autocomplete, not a lookup tool**
- C Connecting to the internet and retrieving live answers
- D Asking expert systems for verified responses

2

IDEA 1 — PREDICTION OVER RETRIEVAL

A student asks AI about a very rare local festival in his village that has never been written about online. The AI gives a confident, detailed answer. What most likely happened?

- A AI found the information from a local government database
- B AI used deep research mode to find obscure sources
- C AI fabricated a plausible-sounding answer — where training text is sparse, prediction produces confident fabrications**
- D AI correctly retrieved facts from its reasoning engine

3

IDEA 1 — MENTAL MODEL

The module says the correct mental model for AI is NOT a "librarian retrieving facts" — but instead a:

- A Search engine crawling the web
- B Expert system with verified rules
- C Writer continuing a narrative — predicting what text most plausibly comes next**
- D Calculator solving logical problems

4

IDEA 2 — ONE-TIME LEARNING, PERMANENT FREEZING

Fatima corrects an AI mistake during a chat: "No, that date is wrong — the event happened in 1947, not 1857." She then starts a brand new chat and asks the same question. What will happen?

- A AI will make the same mistake — corrections during chat don't change the model's weights; no actual learning occurs**
- B AI will remember the correction because memory features updated the model

- C** AI will give a better answer because it re-trained on the feedback
- D** AI will ask for clarification before answering

5

IDEA 2 — KNOWLEDGE CUTOFF

A journalist asks AI: "What happened in the government budget announced last month?" The AI gives an answer as if it knows. Why is this dangerous to trust?

- A** AI is biased against government news
- B** AI's web search mode was not activated
- C** Information after the training cutoff simply doesn't exist in AI's weights — it cannot know recent events
- D** The journalist should have asked in English instead of Urdu

6

IDEA 2 — MEMORY FEATURES

Claude's "Memory" feature remembers that Imran prefers formal writing. Does this mean Claude has now "learned" and updated its AI brain (weights)?

- A** Yes — memory features permanently update the model's weights
- B** No — memory saves facts as text and reinserts them into the context window; weights never change
- C** Yes — this is how AI trains on user feedback
- D** No — memory features only work in paid plans and do not affect anything

7

IDEA 3 — NO INTERNAL TRUTH VERIFICATION

AI confidently states: "According to the Journal of Medicine (2023), aspirin cures cancer." A doctor checks and finds this citation does not exist. This is called hallucination. According to Module 3, why does this happen?

- A** The AI is deliberately lying to seem more intelligent
- B** The AI's internet connection retrieved a wrong source
- C** AI has no internal truth-checker — it only generates; no separate process audits predictions for accuracy before delivery
- D** The AI model was outdated and needed updating

8

IDEA 3 — HALLUCINATION

According to the module, hallucination is NOT a malfunction. What is it?

- A** A bug that developers need to fix
- B** A result of AI accessing wrong websites
- C** The machine working as designed — the same mechanism produces correct answers AND fabrications with identical fluency
- D** A security vulnerability in the AI system

PART 2 BEHAVIORAL EXPLANATIONS — FOUR SUPPORTING IDEAS (CONCEPTS 4–8)

9

IDEA 4 — TOKEN-BASED PROCESSING

Someone asks AI: "How many letter R's are in the word 'strawberry'?" The AI says "2" but the correct answer is 3. Why does AI make this kind of mistake?

- A AI doesn't know the English alphabet
- B AI processes tokens (word chunks), not individual letters — letters inside tokens are essentially invisible to AI
- C AI deliberately gives wrong answers to test users
- D AI's reasoning mode was turned off

10

IDEA 4 — TOKENS & LANGUAGE

A company in Pakistan wants to use AI for Urdu-language customer service. Their developer warns that costs may be higher and context window effectively smaller than for English. Why?

- A AI companies charge more for Urdu language support
- B Urdu requires more server processing time
- C Non-English languages tokenize into more chunks per word due to English-dominant training data, increasing costs and reducing effective context
- D Urdu AI models are less intelligent than English models

11

IDEA 4 — TOKENS & TASKS

Ali asks AI to find a word that rhymes with "orange" and creates an anagram of "listen." AI struggles with both. According to Module 3, why?

- A These tasks require internet access which AI doesn't have
- B Rhyming and anagrams operate on letter-level patterns, but AI processes tokens — individual letters inside tokens are invisible

- C** AI is not trained on poetry or word games
- D** AI needs the reasoning mode enabled for such tasks

12

IDEA 5 — CONTEXT WINDOW AS SOLE INFORMATION SOURCE

Hina started a long chat with AI about her project. Near the end, AI seems to "forget" details from the beginning of the conversation. What is the BEST explanation?

- A** AI got tired and reduced its performance
- B** Hina's internet connection degraded
- C** The context window became overcrowded — older content diluted or summarized away, leaving it outside readable memory
- D** AI deliberately ignores old messages to save processing power

13

IDEA 5 — CONTEXT WINDOW

The module uses a powerful analogy for the context window. It says the context window is like a:

- A** Human brain that stores all past memories
- B** Hard drive that permanently saves all information
- C** Reading desk — information on the desk gets read carefully; information off the desk remains invisible
- D** Cloud server that backs up all conversations

14 **IDEA 5 — BRIEFING PRINCIPLE**

Someone says "I'm being polite when I give AI lots of context and background information." According to Module 3, why is this wrong framing?

- A** Giving too much context slows AI down
- B** AI works better with short prompts
- C** Briefing isn't polite — it's placing information into the ONLY location AI can read; it's a technical necessity, not courtesy
- D** Context makes AI sycophantic

15 **IDEA 6 — CONFIDENCE AS LEARNED STYLE**

An AI states: "Studies show that 73.4% of employees prefer remote work" in a very assured tone. The user cannot find this statistic anywhere. According to Module 3, what happened?

- A** AI accessed a premium research database
- B** AI was trained toward confident, fluent responses — confidence is house style, NOT a signal of truth; it invents statistics with the same tone as real ones
- C** The statistic is real but from a paywalled source
- D** AI's web search mode retrieved an incorrect webpage

16 **IDEA 6 — SYCOPHANCY ORIGIN**

Why does AI tend to agree with users even when users are wrong? According to the module, this "sycophancy" comes from:

- A** AI companies programming AI to be polite

- B** Human feedback training — AI learned that agreeable, helpful, fluent responses get positive ratings, so it trained toward agreement
- C** AI lacking enough training data to disagree
- D** AI's token processing making it unable to form counterarguments

17

IDEA 7 — THE JAGGED FRONTIER

A teacher is amazed that AI brilliantly explains quantum mechanics. She then asks AI to count letters in a simple word and AI gets it wrong. She is confused — how can AI do hard things but fail at easy things? The module explains this as:

- A** AI was tired after the hard question
- B** The AI model was low quality
- C** The Jagged Frontier — AI competence doesn't match human intuition about task difficulty; it's unpredictably brilliant and unpredictably bad
- D** Quantum mechanics questions activate a special reasoning mode

18

IDEA 7 — JAGGED FRONTIER HABITS

Because of the Jagged Frontier, which habit should professionals follow when using AI for important work?

- A** Trust AI on all tasks since it passed the hard ones
- B** Only use AI for simple tasks to avoid errors
- C** Verify at frontier boundaries; don't assume hard-task success predicts easy-task success; re-test AI capabilities quarterly

D Switch to a different AI tool for every task

19

IDEA 8 — TOOLS ENABLE ACTION

Without any tools connected, AI can only do text prediction. But when tools are wired in (like web search or code execution), what happens?

A AI becomes a different model with new weights

B AI stops predicting and starts retrieving from databases

C AI predicts actions (calling tools), executes them, gets results back into context, and continues predicting — transforming into an agent

D AI requires human approval for every action

20

IDEA 8 — TOOLS ENABLE ACTION

The same AI model works as a simple chatbot AND as a folder-organizing assistant that reads and moves files. What changed between the two scenarios?

A A completely different AI model was used for file organization

B The AI was retrained on file management data

C Tools were wired in — the same predictor becomes an agent when tools enable it to take actions

D The context window was expanded for the file task

21

IDEA 9 — REASONING AS PREDICTION-PLUS

When reasoning mode is enabled, AI "thinks step by step" and shows its working before giving a final answer. According to Module 3, this is actually:

- A A fundamentally new cognitive ability that humans gave to AI
- B AI connecting to a separate logic processor
- C Extended prediction — AI predicts reasoning steps, puts them in context, then predicts a better final answer from that reasoning**
- D AI accessing verified knowledge databases step by step

22

IDEA 9 — REASONING LIMITATION

Khalid uses reasoning mode and assumes AI's step-by-step answers are now 100% accurate. Is he right?

- A Yes — reasoning mode eliminates all hallucinations
- B Yes — showing work means every step is verified
- C No — reasoning models still lack a truth-checker; they catch many errors but still hallucinate confidently within rigorous-looking chains**
- D Yes — but only for math and science questions

23

DEMONSTRATION 1 — PREDICTION WITHOUT LOOKUP

To prove that AI predicts rather than retrieves, the module suggests asking AI to explain the rules of a completely made-up, non-existent game. What should happen?

- A AI will say "I don't know this game"**

B AI will search the internet for the game

C AI will confidently fabricate detailed, plausible-sounding rules — proving it predicts, not retrieves

D AI will ask the user to provide the rules first

24

DEMONSTRATION 2 — LEARNING FAILS TO STICK

To prove that AI doesn't truly learn from corrections, what experiment does Module 3 suggest?

A Ask the same question 10 times in one chat and see if answers improve

B Correct a fact in one chat → start a fresh new chat → ask the same question → AI gives the same wrong answer

C Upload a document with corrections and check if AI updates

D Use two different AI models and compare their answers

25

DEMONSTRATION 3 — MISSING TRUTH-CHECKER

To prove AI has no truth-checker, the module suggests asking AI for citations on a narrow, obscure topic and then:

A Checking if the citations are in English

B Asking AI to rate how confident it is in each citation

C Actually verifying the citations exist — some will be completely invented yet delivered with full confidence

D Comparing with another AI model's citations

26

DEMONSTRATION 4 — JAGGED FRONTIER

To demonstrate the Jagged Frontier, the module recommends:

- A Asking AI 100 math questions and tracking its score
- B Comparing two AI models on the same task
- C Present one genuinely hard task AI does brilliantly AND one seemingly simple task AI gets wrong — to see mismatched competence**
- D Ask AI to grade its own responses

27

DEMONSTRATION 5 — REASONING EFFECT

To see the practical effect of reasoning mode, what does the module suggest doing?

- A Ask AI to solve a math problem with and without the internet
- B Ask a hard question plainly (no thinking), then again with thinking/reasoning enabled — and compare the quality difference**
- C Upload documents and compare with and without file upload
- D Ask the same question in English vs. Urdu

28

CORE INSIGHT — THE DRIVING ANALOGY

The module starts with a driving analogy: "You can drive a car without understanding engines, but understanding mechanics helps when something unexpected happens." How does this apply to using AI?

- A You should study AI engineering before using any AI tool
- B AI is like a car — it only works when fueled with data
- C Understanding how AI works helps you avoid over-trusting or dismissing it when unexpected behaviors occur**
- D AI tools are only for technical experts who understand the mechanics

29

CONTEXT WINDOW — PRACTICAL APPLICATION

A legal team uploads a 300-page contract, 50 pages of notes, tool descriptions, and their system prompt into one AI session. They notice AI's response quality drops toward the end. What concept best explains this?

- A AI gets bored with long documents
- B The model weights degraded during use
- C Context window overcrowding — too much information on the "reading desk" dilutes signal and older content gets summarized away**
- D The AI needs internet access to process legal documents

30

IMAGE TOKENS

AI struggles to count exactly how many people are in a crowded photo, even though it can describe the scene well. According to Module 3, why does fine-detail recognition in images suffer?

- A AI cannot process images at all without special plugins
- B Images are broken into token-sized patches — fine detail recognition suffers for the same reason letter-counting does; small things inside patches are invisible**

C AI image processing requires a paid subscription

D The photo resolution was too low for AI

31

WEIGHTS VS CONTEXT — KEY DISTINCTION

What is the correct distinction between "frozen weights" and "context window"?

A Weights are temporary; context window is permanent

B Frozen weights = general knowledge from training (never changes); Context window = specific, current information for this session (changes each time)

C They are the same thing with different names

D Weights store user preferences; context window stores training data

32

SYSTEM PROMPT

When you use Claude on a company's website, it behaves very differently from Claude on claude.ai — it might refuse certain topics or have a different personality. According to Module 3's context window concept, why?

A The company installed a different version of Claude

B The company loaded an invisible system prompt into the context window — this shapes Claude's personality and behavior before any user message

C Claude's weights are different for each company

D The company's internet is slower so Claude performs differently

33

PREDICTION LOOP — AGENT BEHAVIOR

An AI agent is given a task: "Research competitors and save a report to my folder." The agent searches the web, reads pages, writes a report, and saves it — all automatically. What describes this behavior at the core level?

- A A different AI model designed specifically for research
- B A prediction-action-result loop: AI predicts action → tool executes → result enters context → AI predicts next action**
- C AI weights updated in real time to complete the task
- D A human programmer directing the steps behind the scenes

34

TYPOS & TOKENS

Someone types "plese explain kwantum mechaniks" with multiple spelling errors. The AI still understands and gives a great answer. According to token processing, why?

- A AI has a built-in spell-checker that corrects before processing
- B Misspelled words still map to similar token chunks — typos rarely affect AI understanding because processing happens at the token level**
- C AI guesses the meaning from context clues in the sentence
- D AI always ignores spelling and focuses on intent

35

SYCOPHANCY FIX

To overcome AI sycophancy, Module 3 suggests which of these prompting solutions?

- A Ask AI to answer in a different language to avoid bias
- B Use shorter prompts to prevent AI from agreeing too much
- C Use neutral framing (not conclusion-first questions) and objective rubrics with numerical scores — numbers resist agreeable faking better than adjectives**
- D Switch to a different AI model that doesn't have training bias

36

SINGLE SENTENCE SUMMARY

Module 3 ends with one sentence that captures the entire course. Which of these is the **CORRECT** summary from the module?

- A "AI is a search engine that finds the right answer from billions of web pages."
- B "AI is a database that retrieves facts — but needs human verification for recent events."
- C "It is a prediction machine that learned by reading and has no organ for truth, so it is fluent everywhere, reliable only where the text was thick, and you are the part that checks."**
- D "AI is a reasoning engine that thinks carefully and always gives accurate, verified answers."

THE ONE SENTENCE THAT EXPLAINS EVERYTHING

*"It is a prediction machine that learned by reading and has no organ for truth, so it is fluent everywhere, reliable only where the text was thick, and **you are the part that checks.**"*

ANSWER KEY — ALL 36 QUESTIONS

Q1 B	Q2 C	Q3 C	Q4 A	Q5 C
Q6 B	Q7 C	Q8 C	Q9 B	Q10 C
Q11 B	Q12 C	Q13 C	Q14 C	Q15 B
Q16 B	Q17 C	Q18 C	Q19 C	Q20 C
Q21 C	Q22 C	Q23 C	Q24 B	Q25 C
Q26 C	Q27 B	Q28 C	Q29 C	Q30 B
Q31 B	Q32 B	Q33 B	Q34 B	Q35 C
Q36 C				

Prepared for Imran Hassan | Module 3 — What AI Actually Is (Crash Course)

Source: agentfactory.panaversity.org | June 2026